Samarth Bhatt

Data Engineer — Cloud Data Platforms

+91-7600182518 · samarthbhatt8285@gmail.com
linkedin.com/in/samarthbhatt25
github.com/samarth8285 · Pune, India

EXPERIENCE

Data Engineer — Rearc — Aug 2026 – Present
Pune, India

- Onboarding to cloud data engineering practices and pursuing advanced platform certifications ahead of first client engagement.

Associate Data Engineer — Shell India Markets Pvt. Ltd. — Aug 2023 – Jul 2026
Bengaluru, India

- Owned end-to-end SLA-critical batch pipelines on Azure Databricks (PySpark, Spark SQL) for financial reporting, regulatory submissions, and downstream analytics, using incremental, idempotent frameworks on Delta Lake.
- Optimized Spark SQL workloads through partitioning, join strategies, and query tuning — cutting cloud compute costs by ~$170K annually.
- Orchestrated enterprise batch workflows via Control-M and strengthened reliability with CI/CD pipelines, automated testing, and branch protection via GitHub Actions.
- Led platform modernization migrating 500+ production jobs with zero downtime, maintaining 100% SLA compliance throughout.
- Built automated data validation, reconciliation, and audit-control frameworks; built a centralized metadata platform (Databricks REST APIs, Unity Catalog, RBAC) cutting data discovery time by ~40%.
- Created a reusable engineering accelerator framework (saving 120+ engineering hours) and designed star-schema models with SCD Type 2 dimensions.
- Authored 7 internal technical guides and mentored 10+ engineers on performance tuning and responsible GenAI tooling use; delivered Power BI dashboards for fuel and EV pricing analysis.

PROJECT

NYC Yellow Taxi Data Pipeline

Databricks Lakeflow pipeline processing 209 monthly NYC TLC trip-record files (2009–2026) into a medallion (Bronze/Silver/Gold) architecture, deployed as a Databricks Asset Bundle. Idempotent landing into Unity Catalog Volumes feeds schema-aware Auto Loader ingestion that reconciles 3 incompatible historical layouts into a single canonical Bronze table; Silver-layer cleaning feeds two Gold aggregate tables powering a 2-page AI/BI dashboard — Trip Summary and Price Hikes.

PySpark · Delta Lake · Databricks Asset Bundles · Auto Loader · Unity Catalog · Medallion Architecture · AI/BI Dashboards — github.com/samarth8285/nyc-yellow-taxi-project

SKILLS

PROGRAMMING	Python (PySpark, Pandas), SQL (Spark SQL, Databricks SQL), Query Optimization
BIG DATA	Apache Spark, Delta Lake, Databricks Lakeflow, Medallion/Lakehouse Architecture, Data Warehousing
CLOUD	Microsoft Azure, Azure Databricks, ADLS Gen2, Unity Catalog, Docker
GOVERNANCE	Data Validation, Reconciliation, Audit Controls, RBAC, Star Schema, SCD Type 2
ORCHESTRATION	Control-M, Git, GitHub Actions, Bash, Databricks Asset Bundles, IaC
ANALYTICS / BI	Databricks SQL, Power BI, AI/BI (Lakeview) Dashboards

CERTIFICATIONS & EDUCATION

Databricks Certified Data Engineer Professional	Aug 2026	Databricks Certified Associate Data Engineer	Jan 2024
Databricks Certified Gen AI Engineer Associate	Oct 2026	Databricks Certified Apache Spark Developer	Jun 2024

B.E. Computer Engineering, AI & Machine Learning — NMIMS, 2019–2023

AWARDS

Most Valuable Contribution, internal hackathon — high-impact enterprise data solutions.

Recognized for leading enterprise documentation, delivery, and standardization initiatives.

Senior leadership recognition for large-scale cloud cost optimization initiatives.

Special recognition for high-impact Power BI dashboards adopted across stakeholder groups.